



# Deep Learning (Homework 3)



Due date : 2021/6/4 23:55:00 (Hard Deadline)

## 1 Generative adversarial network (50%)

In this exercise, you will implement a [Deep Convolutional Generative Network \(DCGAN\)](#) [1] to synthesis images by using the provided celebrities' face dataset. You can download dataset [img\\_align\\_celeba.zip](#) which was collected from the origin website [CelebFaces](#).



1. Construct a [DCGAN](#) with GAN objective, you can refer to the [tutorial website](#) provided by [PyTorch](#) for implementation.

$$\max_D \mathcal{L}(D) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} \log D(\mathbf{x}) + \mathbb{E}_{\mathbf{z} \sim p_{\mathbf{z}}} \log(1 - D(G(\mathbf{z})))$$


$$\min_G \mathcal{L}(G) = \mathbb{E}_{\mathbf{z} \sim p_{\mathbf{z}}} \log(1 - D(G(\mathbf{z})))$$

- (a) **Describe** how you [preprocess](#) the dataset (such as resize, crop, rotate and flip) and **explain** why. (5%)
- (b) **Plot** the [learning curves](#) for both generator and discriminator. (15%)
- (c) **Draw** some samples generated from your generator at [different training stages](#). For example, you may show the results when running at 5<sup>th</sup> and final learning iteration. (5%)

training for 5 iteration



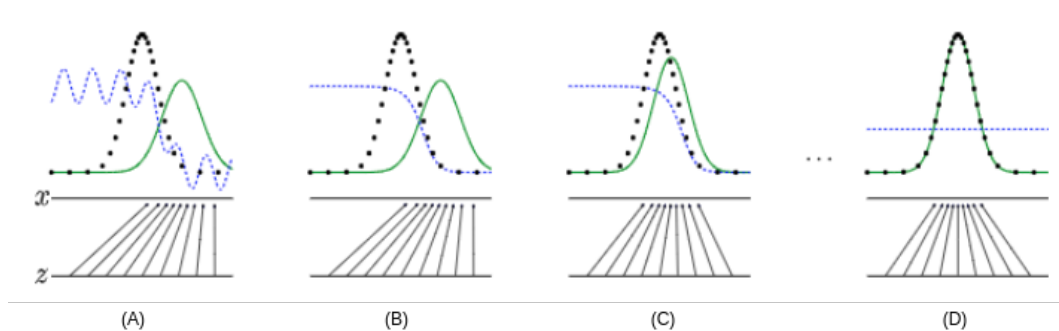
training for 8000 iteration



2. Please **answer the following questions** in your submission report, you can refer to the **paper** to answer these questions. (**Note:** If your answer is more complete and precise, you will receive a higher score.)


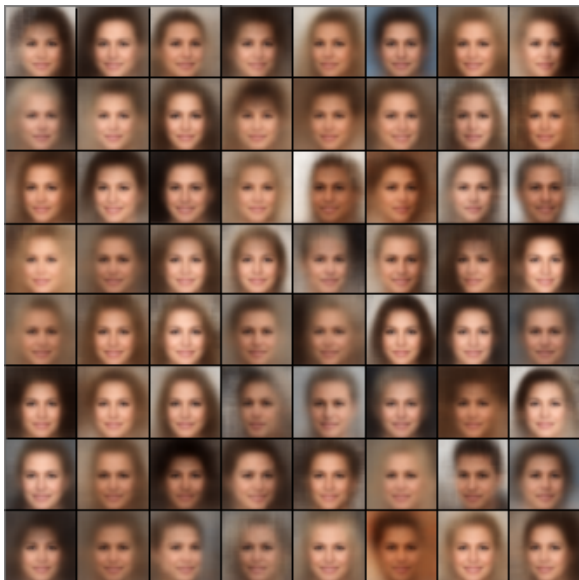
(a) Please **describe** the meaning of the following four pictures during training of GAN, where blue dashed line indicates the **discriminator**, green solid line indicates the **generator**. The answer should include the following: (**Note:** Each step should be discussed.)

- what is the meaning of **black dashed line**,  $x$  and  $z$
- which step is to train the generator or discriminator and show the **corresponding objective function**
- why  $D(x)$  equals to  $\frac{1}{2}$  in ideal case when the training is finished



(b) **The Helvetica Scenario** often happens during training procedure of GAN. Please **explain** why this problem occurs and how to avoid it. (5%)

(c) Both **VAE** and **GAN** are generative models. The following figures are **random generated results** by using VAE (left) and GAN (right). Please compare two results and **describe** the pros and cons of two models. (10%) (**Hint:** You can compare the **loss function** and **training method** using these two models.)




## 2 Deep Q Network (50%)

In this section, you need to clearly know the meaning of **value function** in reinforcement learning. You will calculate the actual **state value**  $V_\pi(s)$  and **state-action value**  $Q_\pi(s, a)$  under different **policies**  $\pi$  for each state-action pair from a given finite Markov Decision Process (MDP). Then you will implement **Deep Q Learning** (DQN) [2] algorithm to approximate the actual  $Q$  value. Finally, you will need to compare the difference between different methods and make some discussions.

$$V_\pi(s) = \mathbb{E}_\pi \left[ \sum_{k=0}^{\infty} \gamma^k R_{t+k+1} \middle| S_t = s \right] = \sum_a \pi(a | s) \sum_{s', r} p(s', r | s, a) [r + \gamma V_\pi(s')]$$

$$Q_\pi(s, a) = \mathbb{E}_\pi \left[ \sum_{k=0}^{\infty} \gamma^k R_{t+k+1} \middle| S_t = s, A_t = a \right] = \sum_{s'} \sum_r p(s', r | s, a) \left[ r + \gamma \sum_{a'} \pi(a' | s') Q_\pi(s', a') \right]$$

1. Please calculate the **state value**  $V_\pi(s)$  and the **state-action value**  $Q_\pi(s, a)$  for state and state-action pair from the following finite MDP under the given policy  $\pi$  and the discount factor  $\gamma = 0.9$ :



- (a) Agent will use the **uniform policy**, which means that  $a_1$  and  $a_2$  will be evenly taken by the agent (8%)
- (b) Agent will use the specific policy which is denoted by  $\pi = [P_{S0}, P_{S1}, P_{S2}, P_{S3}, P_{S4}] = [0, 0, 1, 0, 1]$  where each element represents the probability that the agent will take the action  $a_1$ , in the other word, one minus that probability denotes the probability that agent will take action  $a_2$  (8%)

$a_1$

state	left	right
S0	0.2	0.8
S1	0.7	0.3
S2	0.3	0.7
S3	0	1
S4	0.2	0.8

$a_2$

state	left	right
S0	0.9	0.1
S1	0.2	0.8
S2	0.7	0.3
S3	0	1
S4	0.8	0.2

**Note:** For each policy, you need to calculate a  $8 \times 3$  array to describe the value and  $Q$  value table. Then save the array into .npz file. For example, for the policy in (a) you need to name the file as [value.a.npz](#)

- In this part, you need to implement the Deep  $Q$  Network algorithm to estimate the  $Q$  value from the previous finite MDP by using the DQN agent to interact with the environment.

**Note:** You don't need to implement the environment by yourself. The environment file [finite\\_MDP.env.py](#) is provided by TA. In this environment, the state is defined as a one-hot vector. And the index mapping of each state is [S0, S1, S2, T1, S3, S4, T2, T3] = [0, 1, 2, 3, 4, 5, 6, 7].

- Please follow the algorithm shown below to implement the DQN (10%)
- Please analyze the difference between your answers by the calculation and the output of the DQN (8%)

---

**Algorithm 1** Deep Q-learning with Experience Replay

---

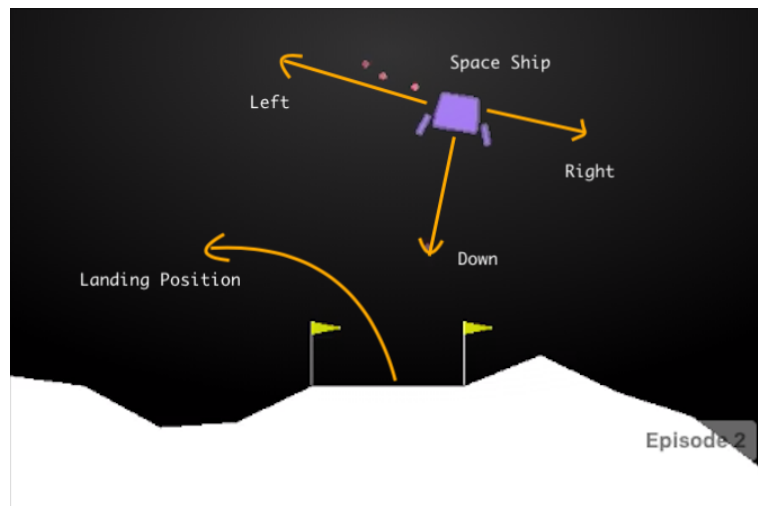
```
Initialize replay memory  $\mathcal{D}$  to capacity  $N$ 
Initialize action-value function  $Q$  with random weights
for episode = 1,  $M$  do
    Initialise sequence  $s_1 = \{x_1\}$  and preprocessed sequenced  $\phi_1 = \phi(s_1)$ 
    for  $t = 1, T$  do
        With probability  $\epsilon$  select a random action  $a_t$ 
        otherwise select  $a_t = \max_a Q^*(\phi(s_t), a; \theta)$ 
        Execute action  $a_t$  in emulator and observe reward  $r_t$  and image  $x_{t+1}$ 
        Set  $s_{t+1} = s_t, a_t, x_{t+1}$  and preprocess  $\phi_{t+1} = \phi(s_{t+1})$ 
        Store transition  $(\phi_t, a_t, r_t, \phi_{t+1})$  in  $\mathcal{D}$ 
        Sample random minibatch of transitions  $(\phi_j, a_j, r_j, \phi_{j+1})$  from  $\mathcal{D}$ 
        Set  $y_j = \begin{cases} r_j & \text{for terminal } \phi_{j+1} \\ r_j + \gamma \max_{a'} Q(\phi_{j+1}, a'; \theta) & \text{for non-terminal } \phi_{j+1} \end{cases}$ 
        Perform a gradient descent step on  $(y_j - Q(\phi_j, a_j; \theta))^2$  according to equation 3
    end for
end for
```

---

**Note:** In the implementation, you can check the provided format of DQN [here](#) to write your own program.

- After the practice of DQN, now we can modify the input and the output sizes and the other parameters of the DQN to train on the openAI gym environment [LunarLander-v2](#)





- (a) Please compare the difference between the total return curves of the previous [finite MDP](#) and the [LunarLander-v2](#) environment. (8%)
- (b) Please choose some hyper-parameters and analyze how these hyper-parameters affect the training result. (8%)

### 3 Rule

- In your submission, you need to submit two files. And only the following file format is accepted:
  - **hw3\_<ProblemNumber>\_<StudentID>.ipynb** file which need to contain all the results, codes and reports for each exercise (e.g. **hw3\_1\_0123456.ipynb**).
- Implementation will be graded by
  - Completeness
  - Algorithm correctness
  - Description of model design
  - Discussion and analysis
- Only [Python](#) implementation is acceptable.
- **DO NOT PLAGIARISM.** (We will check program similarity score.)

### References

- [1] Alec Radford, Luke Metz, and Soumith Chintala, “Unsupervised representation learning with deep convolutional generative adversarial networks,” in *Proc. of International Conference on Learning Representations*, 2016.
- [2] Volodymyr Mnih, Koray Kavukcuoglu, David Silver, Alex Graves, Ioannis Antonoglou, Daan Wierstra, and Martin A. Riedmiller, “Playing Atari with deep reinforcement learning,” *CoRR*, vol. abs/1312.5602, 2013.